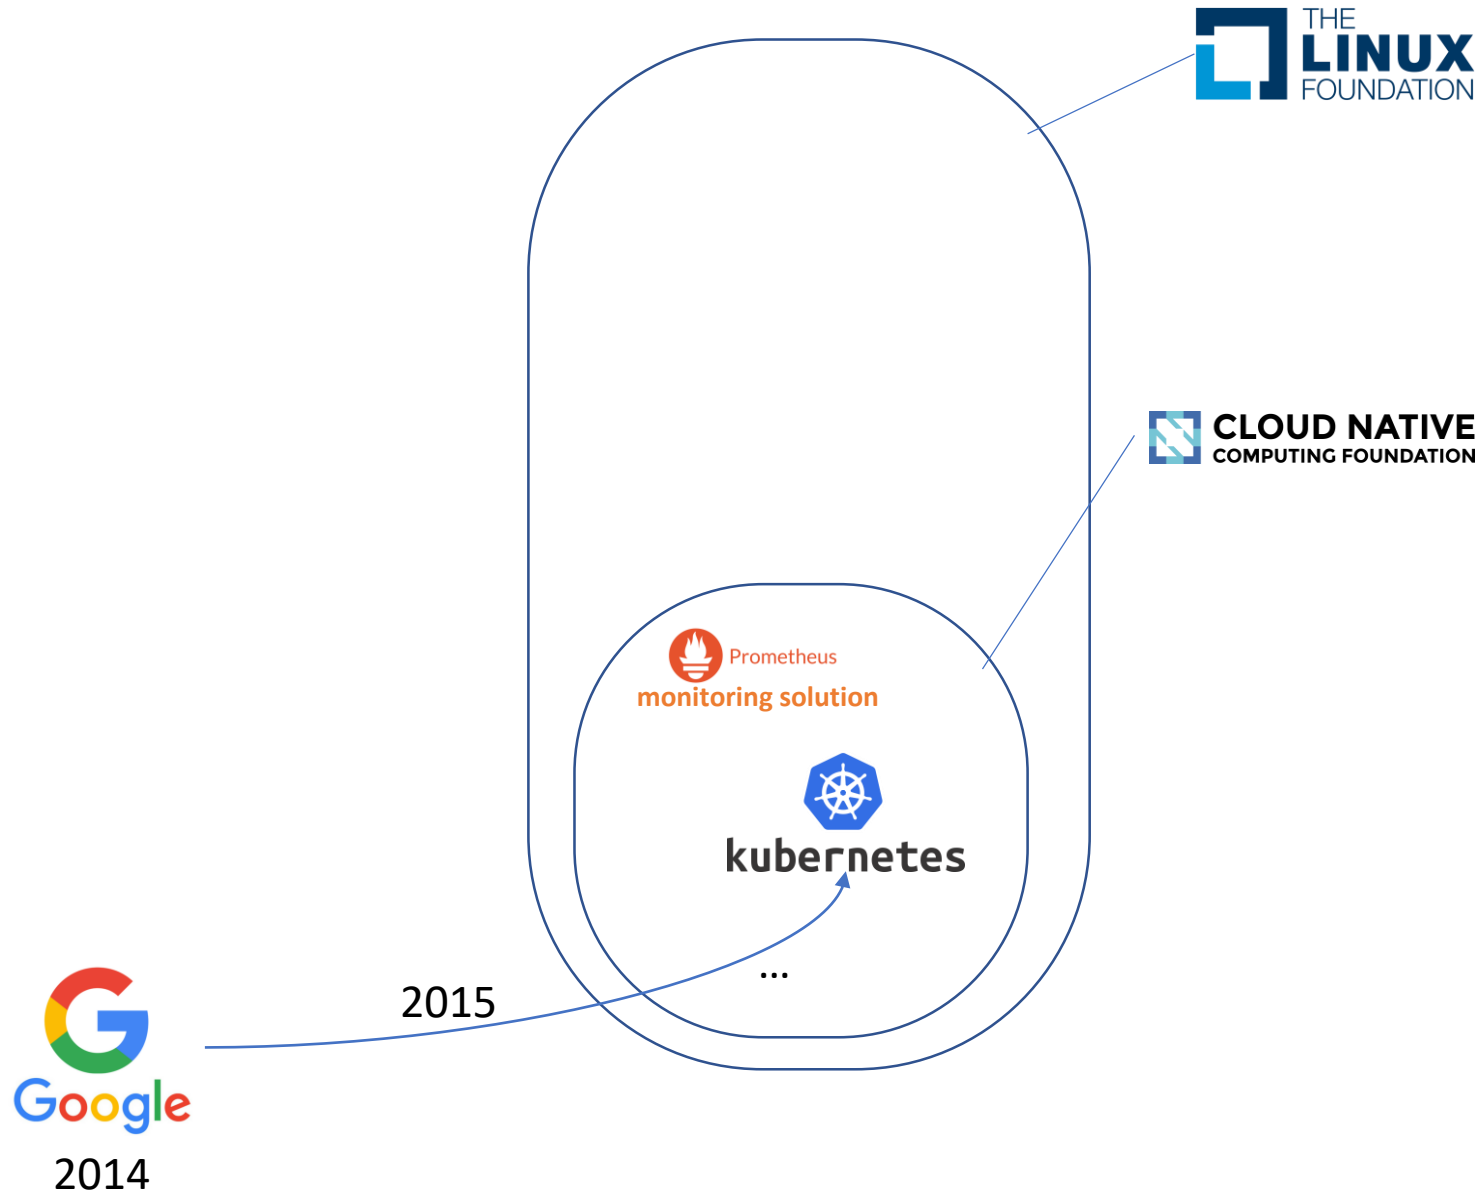


Software Application Engineering

Kubernetes Tutorial

Open-source for the cloud



The Linux Foundation is dedicated to building sustainable ecosystems around open source projects to accelerate technology development and industry adoption [...] through financial and intellectual resources, infrastructure, services, events, and training.

Cloud native technologies empower organizations to build and run scalable applications in modern, dynamic environments such as [...] clouds. Containers, service meshes, microservices, immutable infrastructure, and declarative APIs exemplify this approach.


These techniques enable loosely coupled systems that are resilient, manageable, and observable. Combined with robust automation, they allow engineers to make high-impact changes frequently and predictably with minimal toil.

The Cloud Native Computing Foundation seeks to drive adoption of this paradigm by fostering and sustaining an ecosystem of open source, vendor-neutral projects. We democratize state-of-the-art patterns to make these innovations accessible for everyone.


Container orchestration

In a nutshell:


1. write a deployment YAML file that specifies the configuration to run (“desired state”):
 - containers, exposed ports...
 - starting rank...
 - number of replicas
2. submit it to the Kubernetes Cluster Service
3. This service is in charge of managing the cluster of worker nodes to reach and keep the desired state:
 - a worker node hosts containers
 - If a worker fails, the KCS is responsible for detecting the loss, and reschedule the container(s) somewhere else so that the running configuration matches the desired state




Kubernetes cluster



Kubernetes Cluster





Architecture

Kubernetes Master

Control plane:

- Scheduler
- Network manager

Worker Node 3

Kubelet process

pod 2 replica 1




proxy

Worker Node 1


Kubelet agent

pod 1 replica 1



@IP X

pod 3 replica 1



@IP Y

network proxy

volume
(persistent storage)

Pod: deployment unit


ephemeral

container


Worker Node 2

Kubelet process

pod 1 replica 2




pod 2 replica 2



@IP Z

proxy

Service: set of pods
identified by an IP@ that
does not change when pods
are re-located



Tutorial Series



minikube



Slight differences when running minikube within the VM:

- use driver “none”
- sudo most commands
- ssh not possible



Kubernetes workers run on Ic122 to ic133 only

- Tuto#1: Hello World
 - Discover the Minikube Kubernetes cluster
 - Discover the CLI: your first **kubectl** commands
 - Run the hello-world Docker image in a pod on a node of Minikube
 - Work with a running pod
 - Transfer files between host and containers with `kubectl cp`
- Tuto#2: on Intercell
 - *namespaces*: one cluster, n isolated environments
- Tuto#3: your first YAML deployment file
- Tuto#4: *hostPath volume* to transfer files between host and pod
- Tuto#5: Networking
 - Within a pod
 - Between pods on the same node
 - Between pods on different nodes
 - Service and *NodePort service*

Lab Session (October 20th)

- Deploy your Wikipedia-changes monitor application on InterCell using Docker and Kubernetes
- Demo of kubeview by Patrick Mercier

